



Artificial Intelligence – Lab 04 – Answers

May 10, 2024

Question 1. We can use backward induction, which is equivalent to minimax. Player 1 is playing last, so aims to maximize the value:



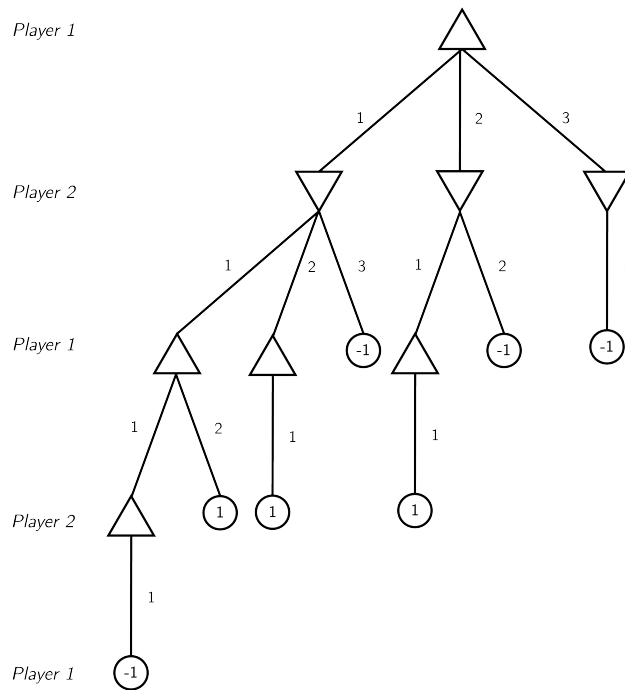
Before that, player 2 was playing, so aiming to minimize the value:



We can now know that, if playing Right first, player 1 is guaranteed to win 5, which is a better payoff than when playing Left.

The optimal sequence is then: [P1: Right, P2: Right, P1: Right]

Question 2. We can represent the task as the following game tree:



where a utility of 1 means that Player 1 (MAX) wins and -1 means that Player 2 (MIN) wins. Using backward induction, we can verify that a **rational** player 2 will always have a solution that guarantees them to win.

Remark: It doesn't mean that Player 1 would never win against any possible Player 2. If Player 2 doesn't play rationally, it may offer the possibility to Player 1 to win. MINIMAX algorithm is guaranteed to be optimal only when the two agents are optimal (i.e. fully rational).

Question 3.

```

1 def minimax_search(state: NimState) -> NimAction:
2     max_utility = -math.inf
3     argmax = None
4     for action in available_actions(state):
5         v = min_value(result(state, action))
6         if v > max_utility:
7             max_utility = v
8             argmax = action
9     return argmax
10
11
12 def max_value(state: NimState) -> int:
13     if is_goal(state):
14         return utility(state)
15     v = -math.inf
16     for action in available_actions(state):
17         v = max(v, min_value(result(state, action)))
18     return v
19
20

```

```

21 def min_value(state: NimState) -> int:
22     if is_goal(state):
23         return utility(state)
24     v = math.inf
25     for action in available_actions(state):
26         v = min(v, max_value(result(state, action)))
27     return v

```

Question 4.

```

1 def maximin_search(state: NimState) -> NimAction:
2     min_utility = math.inf
3     argmin = None
4     for action in available_actions(state):
5         v = max_value(result(state, action))
6         if v < min_utility:
7             min_utility = v
8             argmin = action
9     return argmin

```

One can observe that MIN wins whenever N is a multiple of 4 and that MAX wins otherwise.

Question 5.

- a) Let $N = 4k$ for some $k \in \mathbb{N}$. We proceed by induction over k . The base case was proved in Question 2. Suppose MAX removes $m \in \{1, 2, 3\}$ objects from the pile. A rational MIN can then remove $4 - m$ objects, causing MAX to be faced with $4(k - 1)$ objects. By the induction hypothesis, MAX loses this game against a rational MIN.
- b) If N is not a multiple of 4, it means that $N = 4k + \ell$ for some $k \in \mathbb{N}$ and some $\ell \in \{1, 2, 3\}$. MAX can thus remove ℓ objects from the pile, causing MIN to be faced with $4k$ objects. By a), with the roles of MIN and MAX swapped, we know that MIN must lose against a rational opponent.

Remark: This argument can also be generalized to variants of the game where players are allowed to remove up to $m \in \mathbb{N}$ objects at a time.

Question 6.

```

1 def alpha_beta_search(state: TicTacToeState) -> TicTacToeAction:
2     global count
3     count = 0
4
5     value, action = max_value(state, -math.inf, math.inf)
6     return action
7
8
9 def max_value(
10     state: TicTacToeState, alpha: int, beta: int
11 ) -> tuple[int, TicTacToeAction]:
12     global count
13     count += 1
14
15     if is_goal(state):

```

```

16         return utility(state), None
17     v = -math.inf
18     best_action = None
19
20     for action in available_actions(state):
21         v2, _ = min_value(result(state, action), alpha, beta)
22         if v2 > v:
23             v, best_action = v2, action
24             alpha = max(alpha, v)
25         if v >= beta:
26             break
27     return v, best_action
28
29
30 def min_value(
31     state: TicTacToeState, alpha: int, beta: int
32 ) -> tuple[int, TicTacToeAction]:
33     global count
34     count += 1
35
36     if is_goal(state):
37         return utility(state), None
38     v = math.inf
39     best_action = None
40     for action in available_actions(state):
41         v2, _ = max_value(result(state, action), alpha, beta)
42         if v2 < v:
43             v, best_action = v2, action
44             beta = min(beta, v)
45         if v <= alpha:
46             break
47     return v, best_action

```

Because the AI plays rationally, you can never win against it, but you can end the game in a draw.

Question 7. The number of recursive calls changes, as the order in which actions are explored is shuffled. This is because of the way alpha-beta-search works. Here is an example: Let y and z be two MIN-nodes, which are children of the MAX-node x . Suppose the values of y and z are 1 and 2 respectively (computed by doing minimax on their children). If y is considered before z , all computations will be performed as usual. However, if z is considered before y , we will have $\alpha(x) = 2$, so that, as soon as one of y 's children yields the value 1, the computation will stop: No matter what values the remaining children yield, y 's value will be *at most* 1, and therefore completely irrelevant for x , as it already “knows” it can get a higher value of 2 from z .